

CSA09: DATABASE MANAGEMENT SYSTEMS-ASSIGNMENT QUESTIONS

BY: DHARSHINI J 192324059

Question 1:

ER Diagram Question: Traffic Flow Management System (TFMS)

You are tasked with designing an Entity-Relationship (ER) diagram for a Traffic Flow Management System (TFMS) used in a city to optimize traffic routes, manage intersections, and control traffic signals. The TFMS aims to enhance transportation efficiency by utilizing real-time data from sensors and historical traffic patterns.

Entities and Attributes: (TASK 1)

1.Roads

- Road ID
- Road Name
- Length
- Speed Limit

2.Intersection

- Intersection ID
- Intersection Name
- Latitude
- Longitude

3.Traffic Signals

- Signal ID
- Intersection ID
- Signal status
- Timer

4.Traffic Data

- Traffic Data ID
- Road ID
- Timestamp
- Speed
- Congestion level

Relationship Modelling: (Task 2)

1.Roads to Intersection:

- A road can connect to multiple intersections (one to many).
- An intersection can be connected to multiple roads (many to many).

2. Intersection to Traffic Signals:

- An Intersection can host multiple traffic signals (one to many).
- A Traffic Signal is installed at a single intersection (many to one).

3. Roads to Traffic Data:

- A road can have multiple traffic data (one to many).
- A traffic data record is associated with a single road (many to one).

4. Roads to Traffic Lights:

- A road can have many traffic lights(one to many).
- A traffic light is installed at a single road(one to many).

ER Diagram (Task 3):

Design Justification:

- Scalability
- Real Time Data Processing
- Efficient Traffic Management

Justification and Normalization: (Task 4)

1.First Normal Form (1NF)

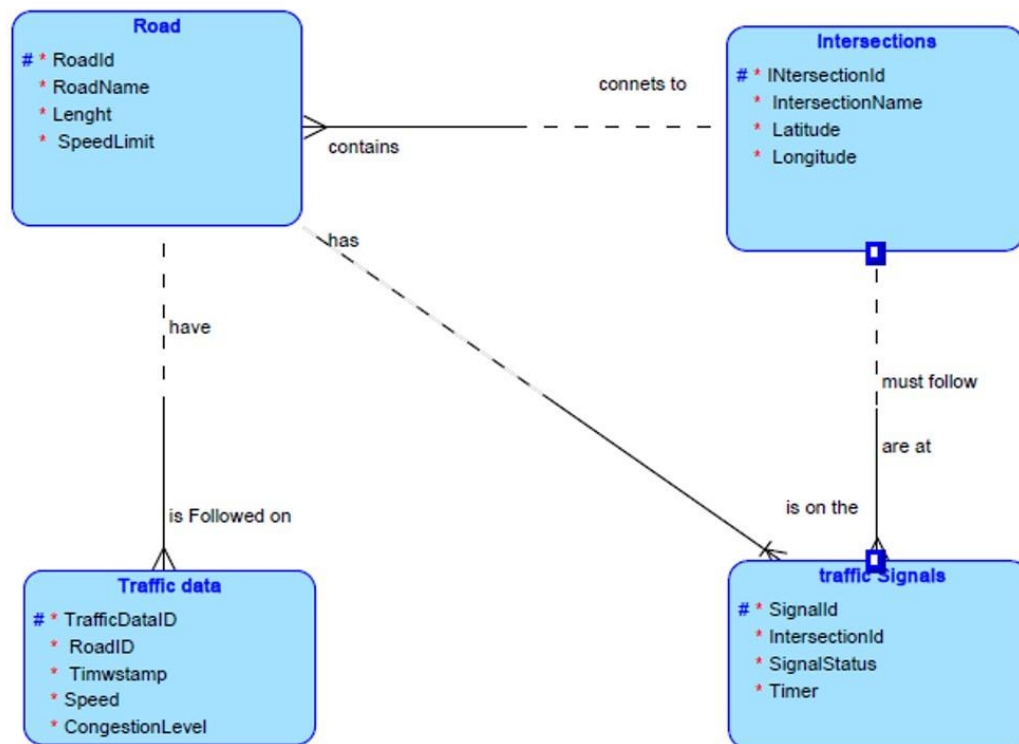
- Ensure each table has atomic columns and there is no repeating groups

2. Second Normal Form (2NF):

- Ensure all non-key attributes are fully functional dependant on the primary key.

3. Third Normal Form(3NF):

- Ensure no transitive dependencies exists between non-key attributes



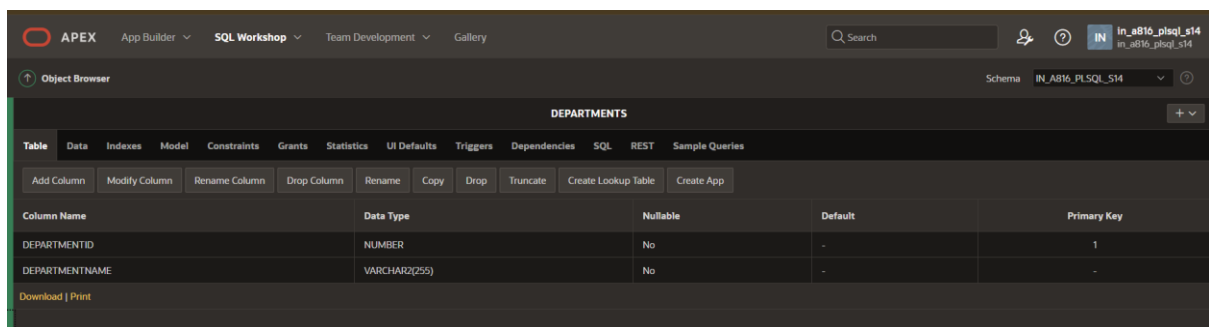
Question 2:

Question 1: Top 3 Departments with Highest Average Salary

Task:

Write a SQL query to find the top 3 departments with the highest average salary of employees. Ensure departments with no employees show an average salary of NULL.

```
CREATE TABLE Departments (  
    DepartmentID INT PRIMARY KEY,  
    DepartmentName VARCHAR(255) NOT NULL  
);
```



The screenshot shows the APEX SQL Workshop interface. The 'Object Browser' tab is active, displaying the 'DEPARTMENTS' table structure. The table has two columns: 'DEPARTMENTID' of type 'NUMBER' and 'DEPARTMENTNAME' of type 'VARCHAR(255)'. Both columns are nullable. The 'DEPARTMENTID' column is marked as the 'Primary Key'.

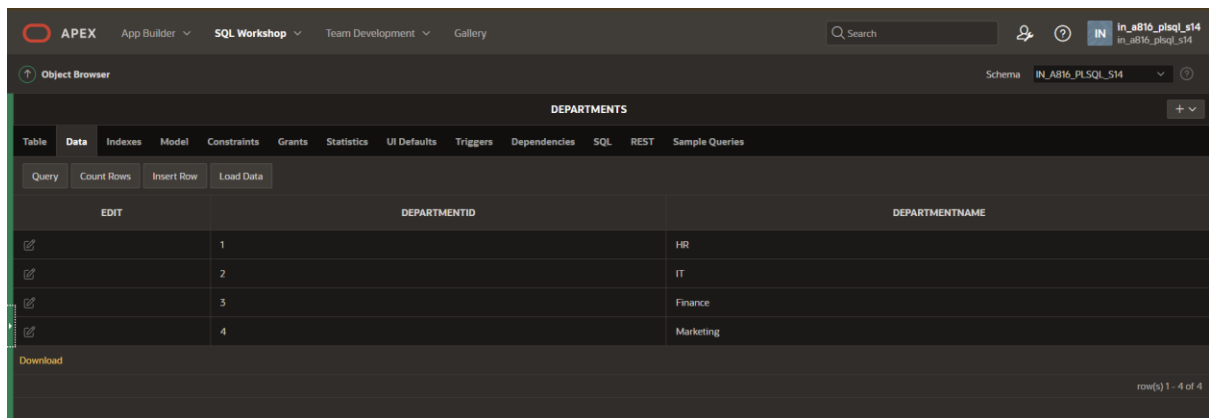
Column Name	Data Type	Nullable	Default	Primary Key
DEPARTMENTID	NUMBER	No	-	1
DEPARTMENTNAME	VARCHAR(255)	No	-	-

```
INSERT INTO Departments VALUES (1, 'HR');
```

```
INSERT INTO Departments VALUES (2, 'IT');
```

```
INSERT INTO Departments VALUES (3, 'Finance');
```

```
INSERT INTO Departments VALUES (4, 'Marketing');
```



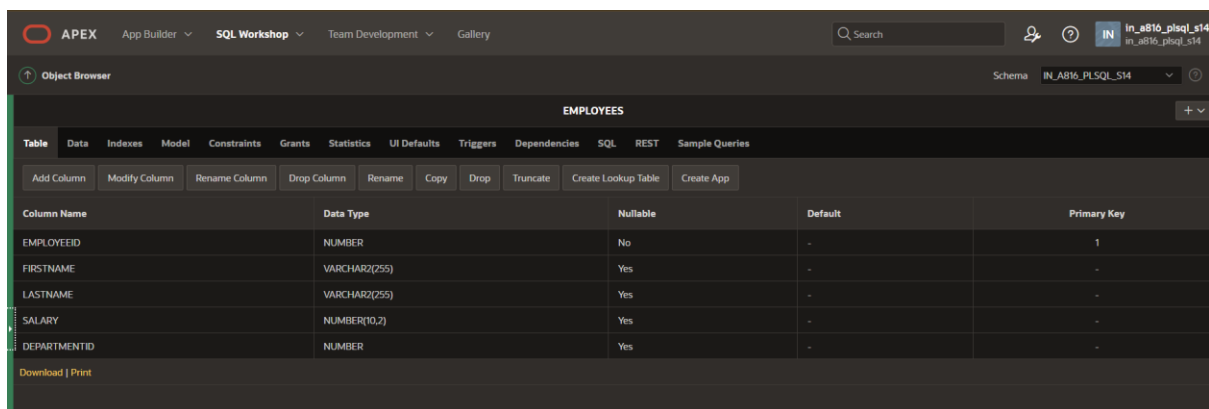
The screenshot shows the APEX SQL Workshop interface. The 'Object Browser' tab is active, displaying the 'DEPARTMENTS' table data. The table has two columns: 'DEPARTMENTID' and 'DEPARTMENTNAME'. The data is as follows:

EDIT	DEPARTMENTID	DEPARTMENTNAME
	1	HR
	2	IT
	3	Finance
	4	Marketing

Download

row(s) 1 - 4 of 4

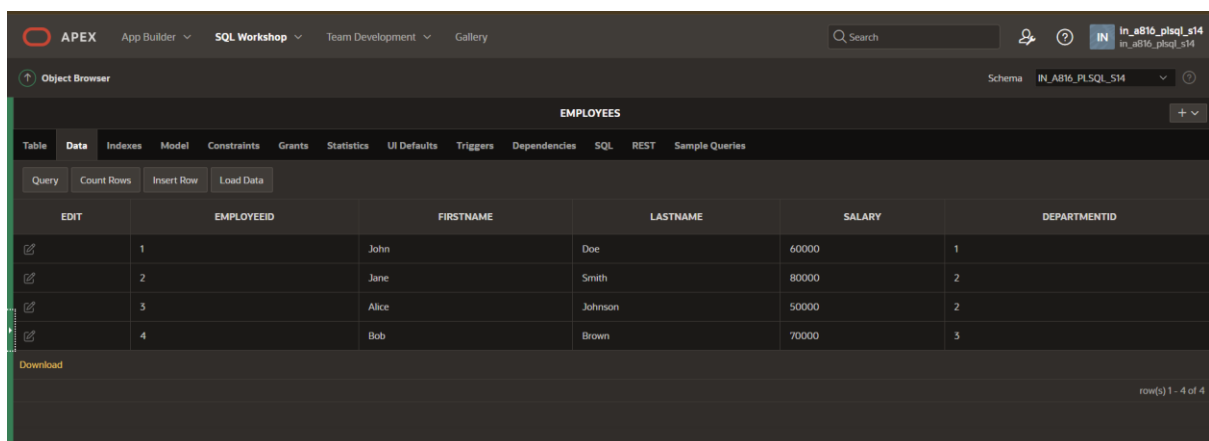
```
CREATE TABLE Employees (
    EmployeeID INT PRIMARY KEY,
    FirstName VARCHAR(255),
    LastName VARCHAR(255),
    Salary DECIMAL(10, 2),
    DepartmentID INT,
    FOREIGN KEY (DepartmentID) REFERENCES Departments(DepartmentID)
);
```



The screenshot shows the Oracle APEX SQL Workshop interface. The 'Object Browser' pane on the left shows the 'EMPLOYEES' table under the 'IN_A816_PLSQL_S14' schema. The 'Table' tab is selected, displaying the table's structure with columns: EMPLOYEEID, FIRSTNAME, LASTNAME, SALARY, and DEPARTMENTID. The 'Data' tab is also visible, showing the table's data.

Column Name	Data Type	Nullable	Default	Primary Key
EMPLOYEEID	NUMBER	No	-	1
FIRSTNAME	VARCHAR(255)	Yes	-	-
LASTNAME	VARCHAR(255)	Yes	-	-
SALARY	NUMBER(10,2)	Yes	-	-
DEPARTMENTID	NUMBER	Yes	-	-

```
INSERT INTO Employees VALUES (1, 'John', 'Doe', 60000, 1);
INSERT INTO Employees VALUES (2, 'Jane', 'Smith', 80000, 2);
INSERT INTO Employees VALUES (3, 'Alice', 'Johnson', 50000, 2);
INSERT INTO Employees VALUES (4, 'Bob', 'Brown', 70000, 3);
```



The screenshot shows the Oracle APEX SQL Workshop interface. The 'Object Browser' pane on the left shows the 'EMPLOYEES' table under the 'IN_A816_PLSQL_S14' schema. The 'Data' tab is selected, displaying the table's data with columns: EMPLOYEEID, FIRSTNAME, LASTNAME, SALARY, and DEPARTMENTID. The data is as follows:

EDIT	EMPLOYEEID	FIRSTNAME	LASTNAME	SALARY	DEPARTMENTID
	1	John	Doe	60000	1
	2	Jane	Smith	80000	2
	3	Alice	Johnson	50000	2
	4	Bob	Brown	70000	3

```
SELECT d.DepartmentID, d.DepartmentName,
ROUND(AVG(e.Salary), 2) AS AvgSalary
```

```

FROM Departments d

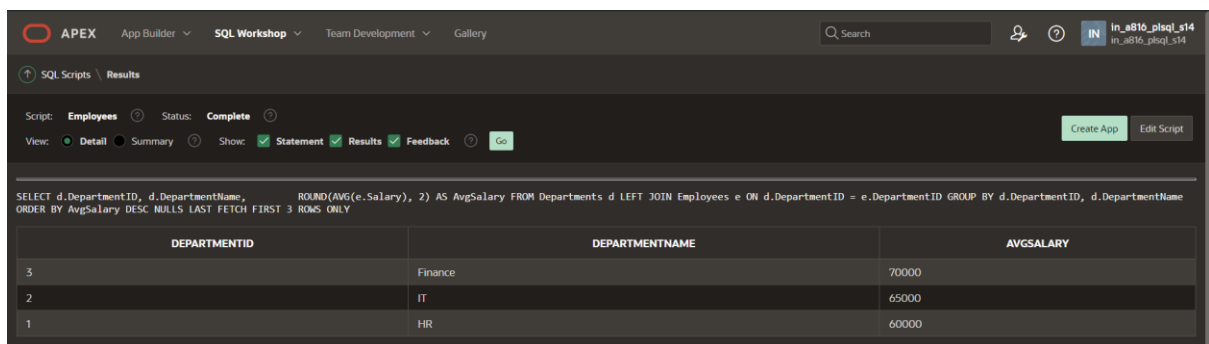
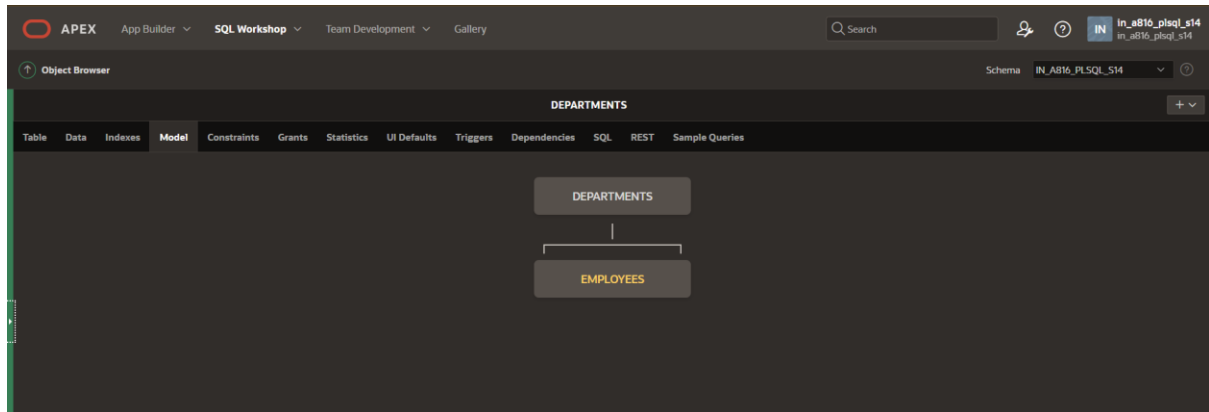
LEFT JOIN Employees e ON d.DepartmentID = e.DepartmentID

GROUP BY d.DepartmentID, d.DepartmentName

ORDER BY AvgSalary DESC

FETCH FIRST 3 ROWS ONLY;

```



Question 2: Retrieving Hierarchical Category Paths

Task:

Write a SQL query using recursive Common Table Expressions (CTE) to retrieve all categories along with their full hierarchical path (e.g., Category > Subcategory > Sub-subcategory).

```

CREATE TABLE Categories (
    CategoryID INT PRIMARY KEY,
    CategoryName VARCHAR(255) NOT NULL,
    ParentCategoryID INT,
    FOREIGN KEY (ParentCategoryID) REFERENCES Categories(CategoryID)
);

```

Object Browser Schema: IN_A816_PLSQL_S14

CATEGORIES

Column Name	Data Type	Nullable	Default	Primary Key
CATEGORYID	NUMBER	No	-	1
CATEGORYNAME	VARCHAR2(255)	No	-	-
PARENTCATEGORYID	NUMBER	Yes	-	-

Download | Print

```
INSERT INTO Categories VALUES (1, 'Electronics', NULL);
```

```
INSERT INTO Categories VALUES (2, 'Computers', 1);
```

```
INSERT INTO Categories VALUES (3, 'Laptops', 2);
```

```
INSERT INTO Categories VALUES (4, 'Smartphones', 1);
```

```
INSERT INTO Categories VALUES (5, 'Furniture', NULL);
```

Object Browser Schema: IN_A816_PLSQL_S14

CATEGORIES

EDIT	CATEGORYID	CATEGORYNAME	PARENTCATEGORYID
	1	Electronics	-
	2	Computers	1
	3	Laptops	2
	4	Smartphones	1
	5	Furniture	-

Download

row(s) 1 : 5 of 5

```
WITH CategoryPath (CategoryID, CategoryName, ParentCategoryID, Path) AS (
```

```
    SELECT CategoryID, CategoryName, ParentCategoryID,
```

```
           CategoryName AS Path
```

```
FROM Categories
```

```
WHERE ParentCategoryID IS NULL
```

```
UNION ALL
```

```
SELECT c.CategoryID, c.CategoryName, c.ParentCategoryID,
```

```
       cp.Path || '>' || c.CategoryName AS Path
```

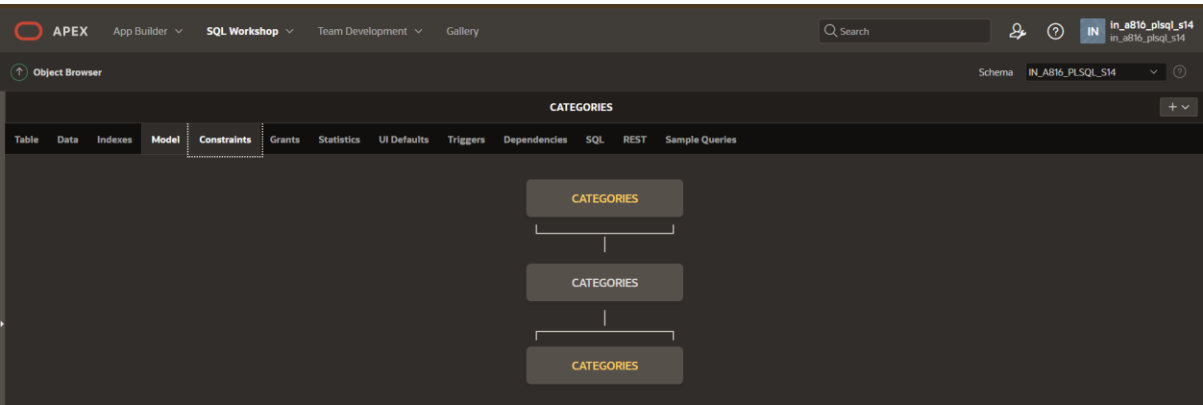
```
FROM Categories c
```

```
JOIN CategoryPath cp ON c.ParentCategoryID = cp.CategoryID
```

```
)
```

```
SELECT CategoryID, CategoryName, Path
```

FROM CategoryPath;



Question 3: Total Distinct Customers by Month

Task:

Design a SQL query to find the total number of distinct customers who made a purchase in each month of the current year. Ensure months with no customer activity show a count of 0.

```
CREATE TABLE Orders
( OrderID INT PRIMARY KEY
, CustomerID INT,
OrderDate DATE,
TotalAmount NUMBER(10, 2) );
```

The screenshot shows the Oracle APEX SQL Workshop interface. The 'Object Browser' on the left shows the 'ORDERS' table. The 'Table' tab is active, displaying a table with columns: ORDERID, CUSTOMERID, ORDERDATE, and TOTALAMOUNT.

Column Name	Data Type	Nullable	Default	Primary Key
ORDERID	NUMBER	No	-	1
CUSTOMERID	NUMBER	Yes	-	-
ORDERDATE	DATE	Yes	-	-
TOTALAMOUNT	NUMBER(10,2)	Yes	-	-

```
INSERT INTO Orders VALUES (1, 101, TO_DATE('2024-01-15', 'YYYY-MM-DD'), 150.00);
INSERT INTO Orders VALUES (2, 102, TO_DATE('2024-02-10', 'YYYY-MM-DD'), 200.00);
INSERT INTO Orders VALUES (3, 103, TO_DATE('2024-02-15', 'YYYY-MM-DD'), 250.00);
INSERT INTO Orders VALUES (4, 101, TO_DATE('2024-03-01', 'YYYY-MM-DD'), 100.00);
```


The screenshot shows the APEX SQL Workshop interface. The 'Object Browser' on the left shows the 'ORDERS' table. The main area displays the table data with columns: ORDERID, CUSTOMERID, ORDERDATE, and TOTALAMOUNT. There are 4 rows of data.

EDIT	ORDERID	CUSTOMERID	ORDERDATE	TOTALAMOUNT
	1	101	15-Jan-2024	150
	2	102	10-Feb-2024	200
	3	103	15-Feb-2024	250
	4	101	01-Mar-2024	100

Download

row(s) 1 - 4 of 4

```

WITH Months AS (
SELECT LEVEL AS MonthNumber
FROM DUAL
CONNECT BY LEVEL <= 12 ),
CustomerCounts AS (
SELECT EXTRACT(MONTH FROM OrderDate) AS MonthNumber,
       COUNT(DISTINCT CustomerID) AS CustomerCount
FROM Orders
WHERE EXTRACT(YEAR FROM OrderDate) = EXTRACT(YEAR FROM SYSDATE)
GROUP BY EXTRACT(MONTH FROM OrderDate) )
SELECT TO_CHAR(TO_DATE(m.MonthNumber, 'MM'), 'Month') AS MonthName,
       COALESCE(c.CustomerCount, 0) AS CustomerCount
FROM Months m
LEFT JOIN CustomerCounts c ON m.MonthNumber = c.MonthNumber
ORDER BY m.MonthNumber;

```

The screenshot shows the APEX SQL Workshop interface. The 'Object Browser' on the left shows the 'CATEGORIES' table. The main area displays the table structure with columns: CATEGORYID, CATEGORYNAME, and PARENTCATEGORYID. The table has 3 columns and 1 primary key (CATEGORYID).

Column Name	Data Type	Nullable	Default	Primary Key
CATEGORYID	NUMBER	No	-	1
CATEGORYNAME	VARCHAR2(255)	No	-	-
PARENTCATEGORYID	NUMBER	Yes	-	-

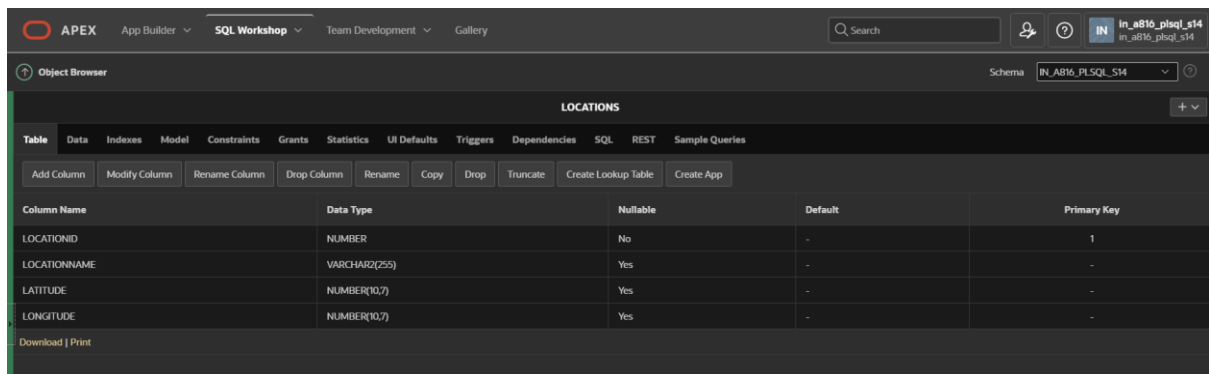
Download | Print

Question 4: Finding Closest Locations

Task:

Write a SQL query to find the closest 5 locations to a given point specified by latitude and longitude. Use spatial functions or advanced mathematical calculations for proximity.

```
CREATE TABLE Locations (  
    LocationID INT PRIMARY KEY,  
    LocationName VARCHAR2(255),  
    Latitude NUMBER(10, 7),  
    Longitude NUMBER(10, 7)  
);
```



The screenshot shows the Oracle APEX SQL Workshop interface. The 'Object Browser' pane on the left shows the 'LOCATIONS' table under the 'IN_A816_PLSQL_S14' schema. The main pane displays the table structure with columns: LOCATIONID (NUMBER, Primary Key), LOCATIONNAME (VARCHAR2(255)), LATITUDE (NUMBER(10,7)), and LONGITUDE (NUMBER(10,7)).

Column Name	Data Type	Nullable	Default	Primary Key
LOCATIONID	NUMBER	No	-	1
LOCATIONNAME	VARCHAR2(255)	Yes	-	-
LATITUDE	NUMBER(10,7)	Yes	-	-
LONGITUDE	NUMBER(10,7)	Yes	-	-

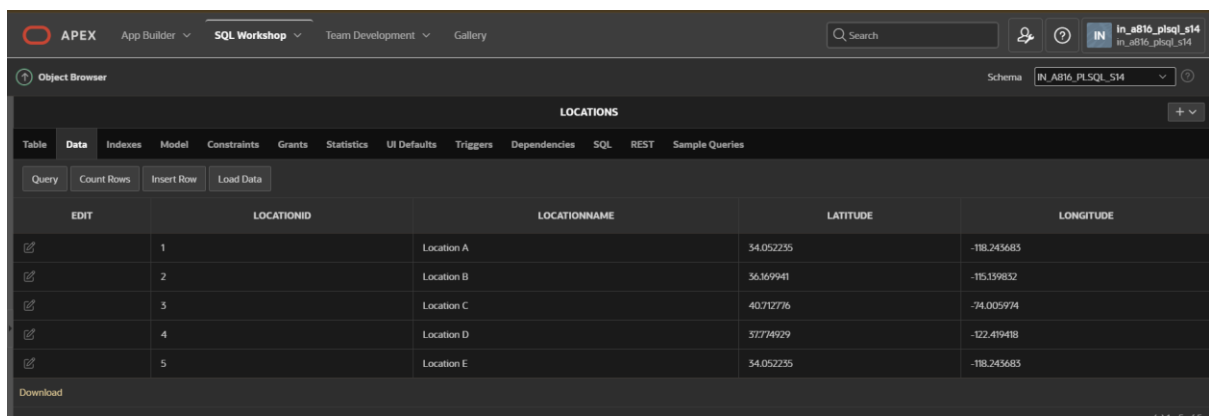
```
INSERT INTO Locations VALUES (1, 'Location A', 34.052235, -118.243683);
```

```
INSERT INTO Locations VALUES (2, 'Location B', 36.169941, -115.139832);
```

```
INSERT INTO Locations VALUES (3, 'Location C', 40.712776, -74.005974);
```

```
INSERT INTO Locations VALUES (4, 'Location D', 37.774929, -122.419418);
```

```
INSERT INTO Locations VALUES (5, 'Location E', 34.052235, -118.243683);
```



The screenshot shows the Oracle APEX SQL Workshop interface. The 'Object Browser' pane on the left shows the 'LOCATIONS' table under the 'IN_A816_PLSQL_S14' schema. The main pane displays the table data with columns: LOCATIONID, LOCATIONNAME, LATITUDE, and LONGITUDE. The data is as follows:

EDIT	LOCATIONID	LOCATIONNAME	LATITUDE	LONGITUDE
	1	Location A	34.052235	-118.243683
	2	Location B	36.169941	-115.139832
	3	Location C	40.712776	-74.005974
	4	Location D	37.774929	-122.419418
	5	Location E	34.052235	-118.243683

```

SELECT LocationID, LocationName, Latitude, Longitude,

        ( 6371 * ACOS(COS((:Latitude * (3.141592653589793 / 180))) * COS((Latitude *
(3.141592653589793 / 180)))

        * COS((Longitude * (3.141592653589793 / 180)) - (:Longitude * (3.141592653589793 / 180)))

        + SIN((:Latitude * (3.141592653589793 / 180))) * SIN((Latitude * (3.141592653589793 / 180))) ) )
AS Distance

FROM Locations

ORDER BY Distance

FETCH FIRST 5 ROWS ONLY;

```

Script: **Locations** Status: **Complete**

View: **Detail** Summary Show: ☒ Statement ☒ Results ☒ Feedback

```

SELECT LocationID, LocationName, Latitude, Longitude,
( 6371 * ACOS(COS((:Latitude * (3.141592653589793 / 180))) * COS((Latitude *
(3.141592653589793 / 180))) - (:Longitude * (3.141592653589793 / 180)))
+ SIN((:Latitude * (3.141592653589793 / 180))) * SIN((Latitude * (3.141592653589793 / 180))) ) ) AS Distance FROM Locations
ORDER BY Distance FETCH FIRST 5 ROWS ONLY

```

LOCATIONID	LOCATIONNAME	LATITUDE	LONGITUDE	DISTANCE
1	Location A	34.052235	-118.243685	-
2	Location B	36.169941	-115.159832	-
5	Location E	34.052235	-118.243685	-
4	Location D	37.74929	-122.419418	-
3	Location C	40.712776	-74.005974	-

5 rows selected. 0.01 seconds

Question 5: Optimizing Query for Orders Table

Task:

- Write a SQL query to retrieve orders placed in the last 7 days from a large Orders table, sorted by order date in descending order.

```

SELECT

OrderID,

CustomerID,

OrderDate,

TotalAmount FROM Orders

WHERE OrderDate >= SYSDATE - INTERVAL '7' DAY

ORDER BY OrderDate DESC;

```

Object Browser

SchemaIN_A816_PL5QL_S14

ORDERS

Table

Data

Indexes

Model

Constraints

Grants

Statistics

UI Defaults

Triggers

Dependencies

SQL

REST

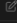
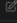
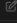
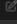
Sample Queries

Query

Count Rows

Insert Row

Load Data

EDIT	ORDERID	CUSTOMERID	ORDERDATE	TOTALAMOUNT
	1	101	15-Jan-2024	150
	2	102	10-Feb-2024	200
	3	103	15-Feb-2024	250
	4	101	01-Mar-2024	100

Download

row(s) 1 - 4 of 4